

# Distributed Sagas

McCaffrey, Caitie                      Kingsbury, Kyle  
Sporty Tights, Inc                      The SF Eagle

Narula, Neha  
That's DOCTOR Narula to you!

May 19, 2015

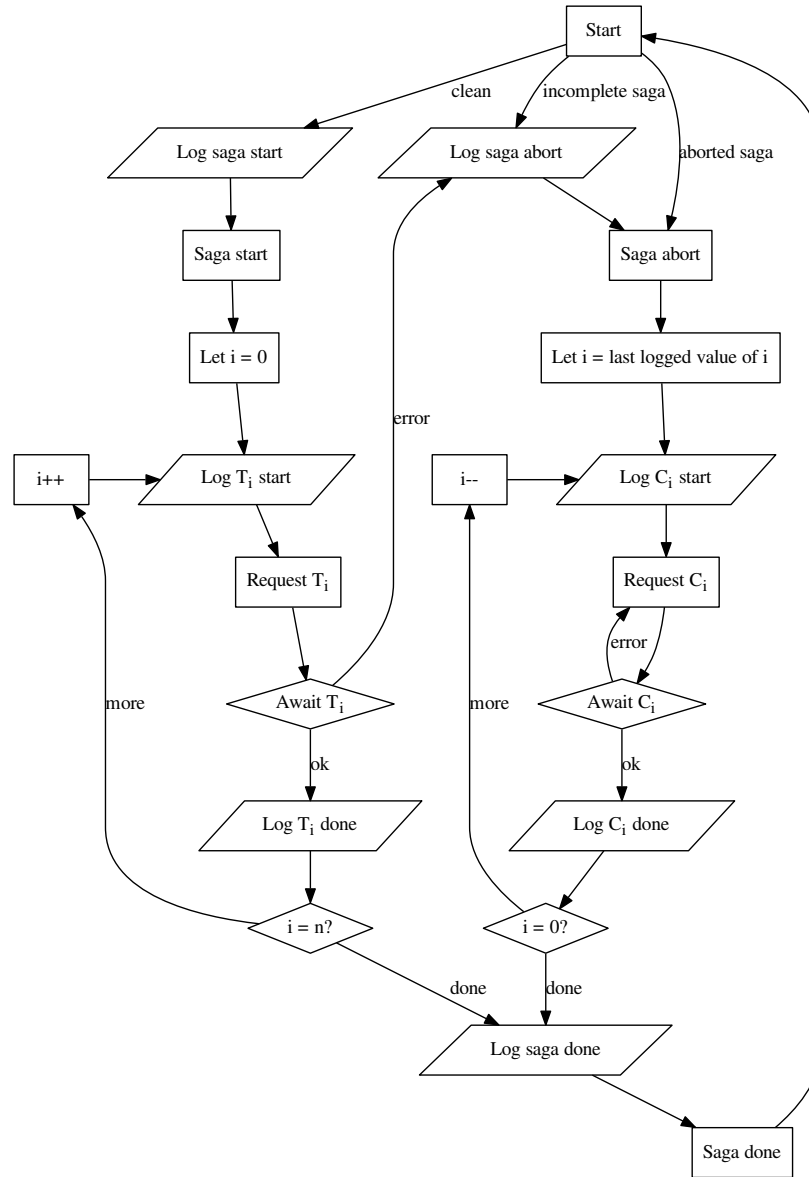
## 1 Introduction

The saga paper outlines a technique for long-lived transactions which provide atomicity and durability without isolation (what about consistency? Preserved outside saga scope, not within, right?). In this work, we generalize sagas to a distributed system, where processes communicate via an asynchronous network, and discover new constraints on saga sub-transactions.

We are especially interested in the problem of writing sagas which interact with *third-party services*, where we control the Saga Execution Coordinator (SEC) and its storage, but not the downstream Transaction Execution Coordinators (TECs) themselves. Communication between the SEC and TEC(s) takes place over an asynchronous network (e.g. TCP) which is allowed to drop, delay, or reorder messages, but not to duplicate them.

We assume a high-availability SEC service running on multiple nodes for fault-tolerance, where multiple SECs may run concurrently. They coordinate their actions through a linearizable data store, which ensures saga transactions proceed sequentially.

## 2 The Saga Execution Coordinator



### 3 Both Rollback and Roll-forward

**Lemma 3.1.** *If  $T_i$  is received by a TEC, then  $T_0, T_1, \dots, T_{i-1}$  have already been acknowledged by a TEC, where  $0 < i \leq n$ .*

*Proof.* In order for  $T_i$  to be received by a TEC, it must have been requested by an SEC. In a roll-forward SEC, this could be a retry of a failed attempt to execute  $T_i$ , but regardless of whether the SEC is roll-back or roll-forward, entering that part of the algorithm requires the SEC to journal its intent to start  $T_i$ .

There are only two paths to that journaling operation. The first case,  $i = 0$ , falls outside our constraint  $0 < i \leq n$ . Therefore the SEC *must* have taken the other path: incrementing  $i$  before beginning a new transaction.

That path depends on  $i - 1 \neq n$  being false, which holds since we are considering  $i \leq n$ . That in turn depends on journaling  $T_{i-1}$ 's completion, which depends on a successful response from a TEC for  $T_{i-1}$ . Therefore some TEC acknowledged  $T_i$ . That in turn requires that TEC to have received  $T_i$ .

So, the receipt of  $T_i$  implies both the receipt and acknowledgement of  $T_{i-1}$ . By induction, receiving  $T_i$  implies *all* transactions  $T_0, T_1, \dots, T_{i-1}$  have been acknowledged. □

**Corollary 3.1.1.** *The first transaction to be received and acknowledged is  $T_0$ .*

*Proof.* Assume the first transaction to be processed is not  $T_0$ , but rather, some  $T_i \mid 0 < i \leq n$ . By 3.1,  $T_{i-1}$  must have been received and acknowledged by a TEC already.  $T_i$  is therefore *not* the first transaction: a contradiction. □

### 4 Rollback

**Lemma 4.1.** *Transactions are requested and received at most once.*

*Proof.* In order for an SEC to request a transaction  $T_i$ , it has to record its intent to execute  $T_i$  in shared SEC storage. Since that storage is linearizable, any other SEC recording an intent to execute  $T_i$  would be visible to the requesting SEC.

**Case 1** Another SEC has already recorded its intent to request  $T_i$ . The given SEC chooses to crash instead of requesting  $T_i$ .

**Case 2** No other SEC has recorded its intent to request  $T_i$ . The given SEC requests  $T_i$  once.

In both cases,  $T_i$  is requested at most once, across all SECs, depending on whether or not the successfully-recording SEC crashes before making its request.

Because the network does not duplicate requests, the number of times  $T_i$  can arrive at a TEC is less than or equal to the number of requests any SEC makes for  $T_i$ . Since that number is at most one,  $T_i$  is received at most once. □

**Lemma 4.2.** *Transactions are seen by TECs in sequential order:  $T_0, T_1, \dots, T_j$ , where  $0 \leq j \leq n$ .*

*Proof.* Consider a sequential history  $S = (T_0, \dots, T_i)$  followed by  $T_j$ . Is  $(T_0, \dots, T_i, T_j)$  sequential? We must show  $i + 1 = j$ .

**Case 1** Assume  $j \leq i$ . Then  $T_j$  is a duplicate of some transaction already in  $S$ , which violates 4.1: a contradiction.

**Case 2** Assume  $i + 1 < j$ . By 3.1,  $T_{i+1}$  must appear before  $T_j$ —but  $S$  cannot contain  $T_{i+1}$ , since it only ranges from 0 to  $T_i$ .

**Case 3** Assume  $i < j \leq i + 1$ . Then  $i + 1 = j$ .

Since cases 1 and 2 are impossible, *any* history comprised of a transaction following a sequential history of at least one element must be sequential as well.

Now, consider histories of one element or fewer:

**Case 1** No transactions occur. The history is trivially sequential.

**Case 2** Exactly one transaction occurs. By 3.1.1, that transaction must be  $T_0$ . This history is trivially sequential.

So any history of one element or fewer is sequential, and any transaction *appended* to that history will also form a sequential history, and so on. By induction, all transactions in a rollback saga system occur sequentially.  $\square$